

ServingMLFastCelery

Example from <https://towardsdatascience.com/deploying-ml-models-in-production-with-fastapi-and-celery-7063e539a5db>

Working example for serving a ML model using FastAPI and Celery.

Usage

Install requirements:

```
pip install -r requirements.txt
```

Set environment variables:

- MODEL_PATH: Path to pickled machine learning model
- BROKER_URI: Message broker to be used by Celery e.g. RabbitMQ
- BACKEND_URI: Celery backend e.g. Redis

```
export MODEL_PATH=...  
export BROKER_URI=...  
export BACKEND_URI=...
```

Start API:

```
uvicorn app:app
```

Start worker node:

```
celery -A celery_task_app:worker worker -l info
```